



It seems like a good idea, but is it scalable?

https://s3.amazonaws.com/lowres.cartoonstock.com/animals-scalable-product-mice-cats-slingshot-mdbn347_low.jpg

COM6012: **Scalable** Machine Learning - Spring 2023

<https://github.com/haipinglu/ScalableML> (Since 2019)

University of Sheffield

Three Instructors + Guest (M. Smith)



Haiping Lu
Module leader



Shuo Zhou



Tahsin Khan

Four Demonstrators (TAs)



Areeb Sherwani (Head)



Pawel Pukowski



Sina Tabakhi



Peizhen Bai

Lectures and Labs

- Lecture: Monday 1pm-2pm @ Diamond LT2
- Lab: Friday 11am-1pm @ Diamond Computer Room 1
 - Finish required part of the labs in this slot!



Other Interactions

- Additional sessions
 - Instructor office hour: Mon 4:30-5:30pm @G25 Regent Court
 - Online Discord help sessions: Wed 2pm-4pm from week 2
- Blackboard discussion board: post your questions to get answered by the course instructors
 - One general forum: general question/feedback
 - Three forums: Haiping, Shuo, Tahsin
 - Get help on lecture/lab contents
 - Ask respective assignment questions
 - Assignment questions:
 - To ask for clarification on assignment questions (i.e. the tasks to do)
 - NOT to ask how to solve the problems, the correctness of a specific solution, or share a possible solution. It is an assessment.
- Direct email to instructors: personal/private issues only

Assessment

- Lab exercises: 0% (self assessment)
 - Finish lab exercises by the following Wednesdays
 - Solutions to release on the following Thursdays
- Assignment: 50%
 - Progressive release → **start EARLY**: complete by 31 March
 - Deadline: 13:00 on Fri, **5th May** (end of lab)
 - Solution release: 12th May
 - Marking and feedback deadline: 26th May
- Final exam: 50%
 - To be scheduled in exam period (22 May – 9 June)
 - Formal exam on Blackboard: 2 hours (with samples available)

Contents: **Very Hands-on**



Week	Date	Topic	Instructor
1	06 Feb	Introduction to Spark and HPC	Haiping
2	13 Feb	RDD, DataFrame, ML pipeline, & parallelization	Haiping
3	20 Feb	Scalable logistic regression	Shuo
4	27 Feb	Scalable generalized linear models	Shuo
5	06 Mar	Scalable matrix fact. for collaborative filtering (RecSys)	Haiping
6	13 Mar	Scalable K-means clustering	Haiping
7	20 Mar	Scalable PCA for dimensionality reduction	Haiping
8	27 Mar	Scalable decision trees	Tahsin
9	24 Apr	<i>Apache spark in the cloud (guest lecture, not assessed)</i>	<i>Mike</i>
12	15 May	<i>Scalable neural networks (not assessed)</i>	<i>Tahsin</i>